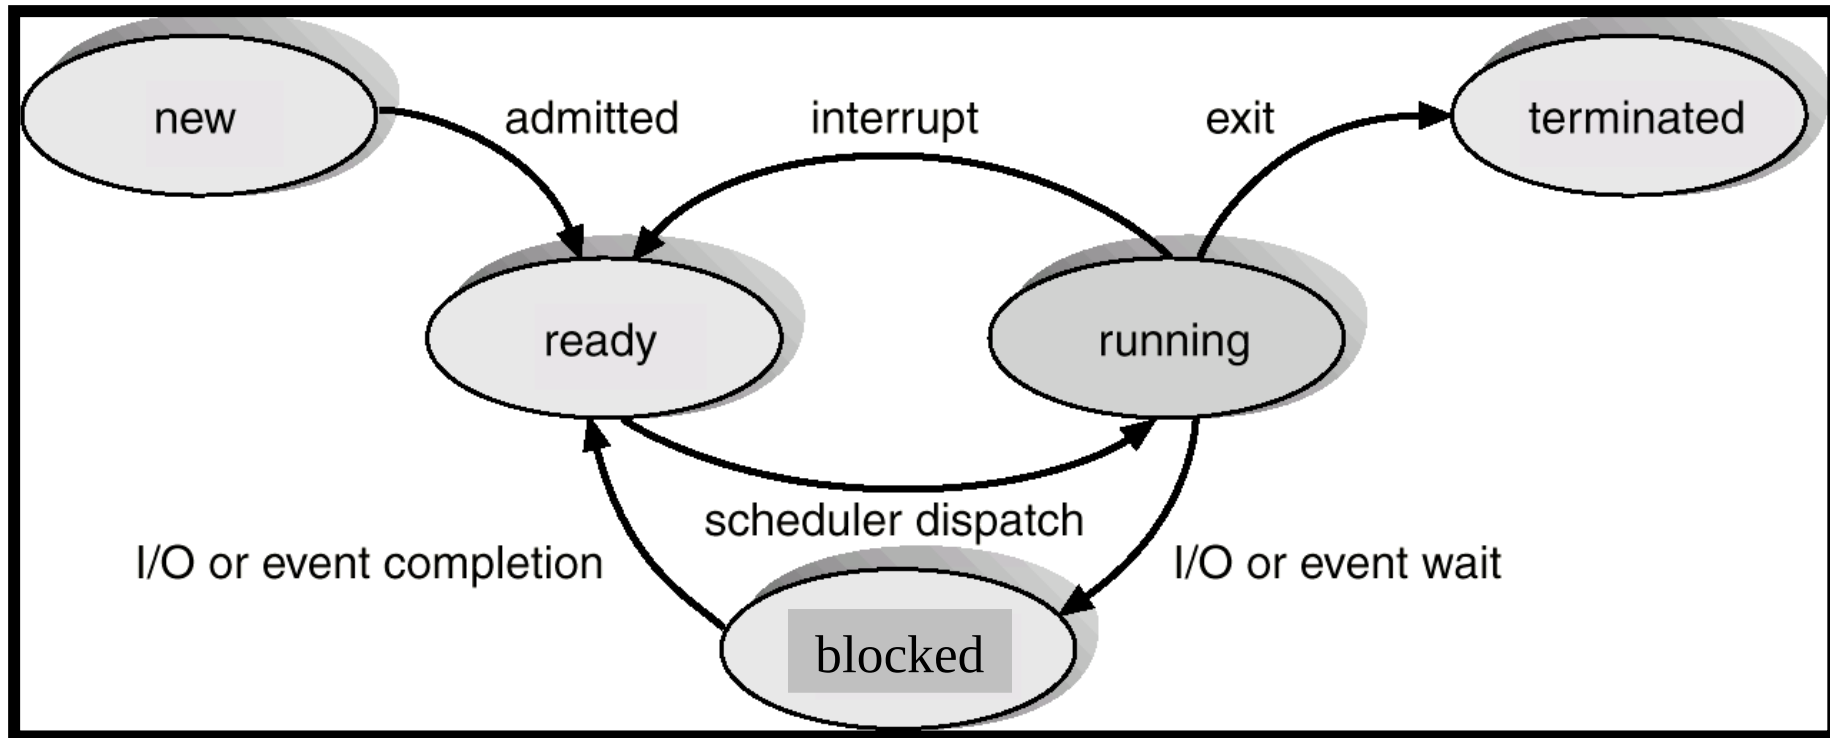


Basic Synchronization Principles

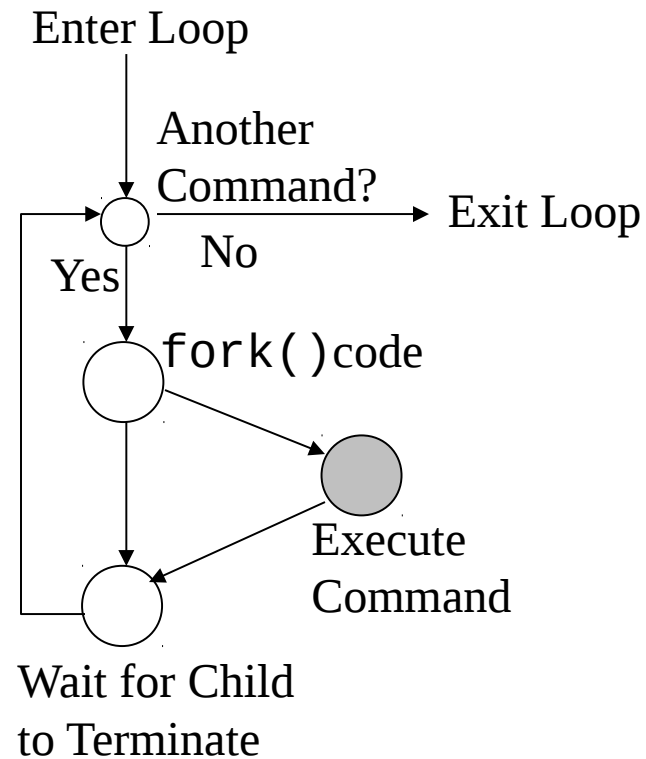
5



Concurrency

- Value of concurrency – speed & economics
- But few widely-accepted concurrent programming languages (Java is an exception)
- Few concurrent programming paradigm
 - Each problem requires careful consideration
 - There is no common model
- OS tools to support concurrency tend to be “low level”

Command Execution



(a) UNIX Shell

Background

- assembly language

$x + y$

```
mov Ax,x  
mov Bx,y  
add Bx
```

if ($x == y$)

```
mov Ax,x  
mov Bx,y  
sub Bx  
jz
```

- a running process can be halted unexpectedly – between two clock cycles, whether or not the execution of an instruction completed
- cannot assume the length, speed, or order of execution
- intermittent bug is next to impossible to replicate
- a bug may not be visible at run-time
- superior programming design on paper
- not adequate to be correct most of the time

critical section problem

example – initial checking balance \$800

// **Teller process**

clear check for \$500

get balance0

if (balance0 > amount)

 deduct from balance0

save balance0

clear next check

// **ATM process**

enter pin

deposit to checking

enter amount \$1000

get balance0

add to balance0

save balance0

next transaction

Critical Sections (2)

- Mutual exclusion: Only one process can be in the critical section at a time
- There is a race to execute critical sections
- The sections may be defined by different code in different processes
 - cannot easily detect with static analysis
- Without mutual exclusion, results of multiple execution are not determinate
- Programming intervention to ensure correctness

Some Possible Solutions

- Disable interrupts
- Software solution – spin locks
- synchronizing objects
- Transactions
- `FORK()`, `JOIN()`, and `QUIT()` [Chapter 2]
 - Terminate processes with `QUIT()` to synchronize
 - Create processes whenever critical section is complete
- ... something new ...

Using a Lock Variable

```
shared boolean lock = FALSE;  
shared double balance;
```

Code for p_1

```
/* Acquire the lock */  
while(lock) ;  
lock = TRUE;  
/* Execute critical sect */  
balance = balance + amount;  
/* Release lock */  
lock = FALSE;
```

Code for p_2

```
/* Acquire the lock */  
while(lock) ;  
lock = TRUE;  
/* Execute critical sect */  
balance = balance - amount;  
/* Release lock */  
lock = FALSE;
```


Solution Constraints

- Mutual exclusion: Only one process at a time in the CS
- Only processes competing for a CS are involved in resolving who enters the CS
- Once a process attempts to enter its CS, it cannot be postponed indefinitely
- After requesting entry, only a bounded number of other processes may enter before the requesting process

Semaphore Assumptions

- Memory read/writes are indivisible (simultaneous attempts result in some arbitrary order of access)
- There is no priority among the processes
- Relative speeds of the processes/processors is unknown
- Processes are cyclic and sequential

Semaphores

A semaphore, sem, is an integer variable that can only be changed or tested by these two indivisible functions:
P (sem) and V (sem)

```
P ( semvar )  
    decrement semvar  
    if ( semvar is negative )  
        block the calling process  
        (place in blocked queue)  
endP
```

```
V ( semvar )  
    increment semvar  
    if ( semvar queue is not empty )  
        resume one process from the queue  
endV
```

Notes:

A process executes a P operation to block itself

A running process resumes another process with the V operation

A resumed process may need to wait for its turn to execute

Bounded buffer problem

- shared buffer that can hold 3 items ($n = 3$)
- producer process puts item (one by one) in an available buffer, i.e., the producer must halt when all 3 spaces are filled
- consumer process takes an available item (one by one, in order) from a buffer, i.e., the consumer must halt when no item is available
- two semaphores: $PUT_ITEM = 3$, $TAKE_ITEM = 0$

```
// producer process  
loop  
    produce an item  
    put an item  
    increment pointer  
endloop
```

```
// consumer process  
loop  
    take one item from buffer  
    increment pointer  
    consumer item  
endloop
```

P(PUT_ITEM) - This operation determines whether or not there is an empty slot in the buffer memory. If there is an empty slot, the producer copies one item to buffer memory. If all slots are full, the OS will block the producer process (moves it to the blocked queue for this semaphore). The semaphore variable PUT_ITEM denotes the state of the buffer - when its value is negative after P decrements PUT_ITEM, all slots are full.

P(TAKE_ITEM) - This operation determines whether or not any slot contains an item to be taken. A negative value of the semaphore variable (as P(TAKE_ITEM) executes) denotes that all slots are empty; therefore, the OS will block the consumer process.

V(TAKE_ITEM) - After one item is copied to buffer memory, the producer needs to alert the consumer process that there is one more item to take (when it gets a chance). If the consumer process happened to be blocked when V(TAKE_ITEM) executes, the OS will resume the consumer process (moves it to ready queue). If the consumer is not blocked, by incrementing the semaphore TAKE_ITEM value, the consumer process will not be blocked when it executes the P(TAKE_ITEM) function.

V(PUT_ITEM) - When one item is taken from buffer memory, the producer process is alerted that there is one more empty slot. The OS will resume the producer process, if it is blocked when V(PUT_ITEM) executes. If it's not blocked, the OS won't block it when the producer executes P(PUT_ITEM) because the value of semaphore PUT_ITEM will not go negative.

Test cases

Some test cases

- 1) Can the producer process fill the buffer?
- 2) Can the consumer process empty the whole buffer?
- 3) Can the producer process put a new item after the buffer is already full?
- 4) Can the consumer process take an item when the buffer is empty?
- 5) Any deadlock? Any starvation?

some hints

1. placing P() operations: does there exist any instance where a process is required to block before this instruction executes?
2. determining the number of semaphores to use by identifying the requirements/reasons for blocking a process
3. setting initial value of each semaphore: use the largest initial value possible and use zero initial value if a process is required to block immediately (assuming that it is the first to execute)
4. placing V() operations: resume a blocked process as soon as possible

critical section example

Initial checking balance \$800

```
// TellerA process  
clear check for $500  
get balance0  
if (balance0 > amount)  
    deduct from  
    balance0  
save balance0  
clear next check
```

```
// ATM process  
enter pin  
deposit to checking  
enter amount $1000  
get balance0  
add to balance0  
save balance0  
next transaction
```

```
// TellerB process  
clear check for $100  
get balance0  
if (balance0 > amount)  
    deduct from  
    balance0  
save balance0  
clear next check
```


synchronizing

- three concurrent processes
- the first process outputs 'A', second outputs 'B', third outputs 'C'
- output pattern: ACBCACBCACBCACBC... (A always first)

Coordination

All processes finish the F functions before starting the G functions.
All functions execute independently.

loop repeatedly

F_one

G_one

H_one

endloop

loop repeatedly

F_two

G_two

H_two

endloop

loop repeatedly

F_three

G_three


H_three

endloop

Device Controller Synchronization

- The semaphore principle is logically used with the **busy** and **done** flags in a controller
- Device driver signals device controller with a **V(busy)**, then waits for completion with **P(done)**
- Controller waits for work with **P(busy)**, then announces completion with **V(done)**
- What are the initial values?

Dining-Philosophers Problem



- each philosopher performs each task sequentially and repeatedly
- three tasks: thinking, waiting to eat, eating
- to eat, a philosopher must pick up two chop sticks – on the left and on the right
- chopsticks represents computing resources: disk space, modem, printer, etc.
- maximum concurrency, no starvation, no deadlock

Dining-Philosophers Problem

THINKING=0, HUNGRY=1, EATING=2,
LEFT=(k-1)%5, RIGHT=(k+1)%5

semaphore mutex // initially 1
semaphore sem[5] // initially 0
int state[5];

take_forks(int k)
 P(mutex)
 state[k]=HUNGRY
 test(k)
 V(mutex)
 P(sem[k])

```
test(int k)
    if (state[k]==HUNGRY &&
        state[LEFT]!=EATING &&
        state[RIGHT]!=EATING)
        state[k]=EATING
        V(sem[k])
```

Dining-Philosophers Problem


```
put_forks(int k)
    P(mutex)
    state[k]=THINKING
    test(LEFT)
    test(RIGHT)
    V(mutex)
```

```
test(int k)
    if (state[k]==HUNGRY &&
        state[LEFT]!=EATING &&
        state[RIGHT]!=EATING)
        state[k]=EATING
    V(sem[k])
```

*this implementation leaves
much to be desired:*

- 1. starvation*
- 2. lacking concurrency*

Readers-Writers Problem (3)



Implementing Semaphores

- Minimize effect on the I/O system
- Processes are only blocked on their own critical sections (not critical sections that they should not care about)
- If disabling interrupts, be sure to bound the time they are disabled

Synchronization Hardware

- atomic operation must execute entirely without interruption
- Test and modify the content of a word atomically

```
boolean TestAndSet (donotenter, temp, true) {  
    donotenter = temp;  
    temp = true;  
    return donotenter;  
}
```

Active vs Passive Semaphores

- A process can dominate the semaphore
 - Performs V operation, but continues to execute
 - Performs another P operation before releasing the CPU
 - Called a passive implementation of V
- Active implementation calls scheduler as part of the V operation.
 - Changes semantics of semaphore!
 - Cause people to rethink solutions